


# Introduction to Supervised Learning

Dr. Qiwei Gan



# A Basic Machine Learning Workflow



# Input Data as a Table



```
In [2]: labeled_images = pd.read_csv('C:\\data\\digits_train.csv')
```

```
In [4]: labeled_images.head()
```

Out[4]:

	label	pixel0	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	...
0	1	0	0	0	0	0	0	0	0	0	...
1	0	0	0	0	0	0	0	0	0	0	...
2	1	0	0	0	0	0	0	0	0	0	...
3	4	0	0	0	0	0	0	0	0	0	...
	0	0	0	0	0	0	0	0	0	0	...

Each row corresponds to a single data instance (sample)

The 'label' column contains the label for each data instance (sample)

Index

These columns contains the features of each data instance (sample)

# Conventions



- Capital X: features matrix
- Lower y: label

```
In [4]: #select images features from the second column to the last column.  
X = labeled_images.iloc[:,1:]  
  
#select the first column which is the label, or the digit.  
y = labeled_images.iloc[:,0]  
  
print('The original input dataset shape is: ', labeled_images.shape)  
print('The X dataset shape is: ',X.shape)  
print('The y dataset shape is: ', y.shape)
```


The original input dataset shape is: (42000, 785)

The X dataset shape is: (42000, 784)

The y dataset shape is: (42000, 1)



# Splitting Input Data into Train/Test



Data used to build the machine learning model, are called the *training data*.

The rest of the data will be used to assess how well the model works; these data are called *test data*.

# Training Models



## Create classifier object

```
In [ ]: from sklearn.neighbors import KNeighborsClassifier  
knn = KNeighborsClassifier(n_neighbors = 5)
```

Tuning the parameters

## Train the classifier (fit the estimator) using the training data

```
In [ ]: knn.fit(X_train, y_train)
```

## Estimate the accuracy of the classifier on future data, using the test data

```
In [ ]: knn.score(X_test, y_test)
```

## Use the trained classifier model to classify (predict) new, previously unseen data

```
In [ ]: predictd_y=knn.predict(y_test)
```

# Questions?



# Overfitting and Underfitting



	Classifier	Score on training data	Score on test data
0	Logistic Regression Classifier	1.000000	0.835000
1	Support Vector Machine	1.000000	0.120000
2	Decision Tree Classifier	1.000000	0.680000
3	Random Forest classifier	0.998750	0.847500
4	Ridge Classifier	0.973750	0.745000
5	K-Nearest Neighbors Classifier	0.930000	0.870000
6	SGD Classifier	0.903750	0.835000
7	Naive Bayes Multinomial Classifier	0.843750	0.805000
8	ElasticNet Classifier	0.663894	0.531748
9	Lasso Classifier	0.639250	0.535359

	Number of components	Classifier	Score on training data	Score on test data
0	50	Support Vector Machine	0.994375	0.920000
1	90	Support Vector Machine	0.998125	0.920000
2	30	Support Vector Machine	0.985000	0.915000
3	70	Support Vector Machine	0.996250	0.915000
4	110	Support Vector Machine	0.998125	0.910000
5	130	Support Vector Machine	0.998125	0.905000
6	150	Support Vector Machine	0.998125	0.902500
7	170	Support Vector Machine	0.998750	0.895000
8	190	Support Vector Machine	0.998750	0.887500
9	10	Support Vector Machine	0.928750	0.882500



# Overfitting and Underfitting



- **Generalization** ability refers to an algorithm's ability to give accurate predictions for **new, previously unseen** data.
- **Assumptions:** *Future unseen data (test set) will have the same properties as the current training sets.*
  - *Thus, models that are accurate on the training set are expected to be accurate on the test set.*
  - *But that may not happen if the trained model is tuned too specifically to the training set.*
- Models that are **too complex** for the amount of training data available are said to **overfit** and are not likely to generalize well to new examples.
- Models that are **too simple**, that don't even do well on the training data, are said to **underfit** and also not likely to generalize well.

# Trade-off between Complexity and Generalization


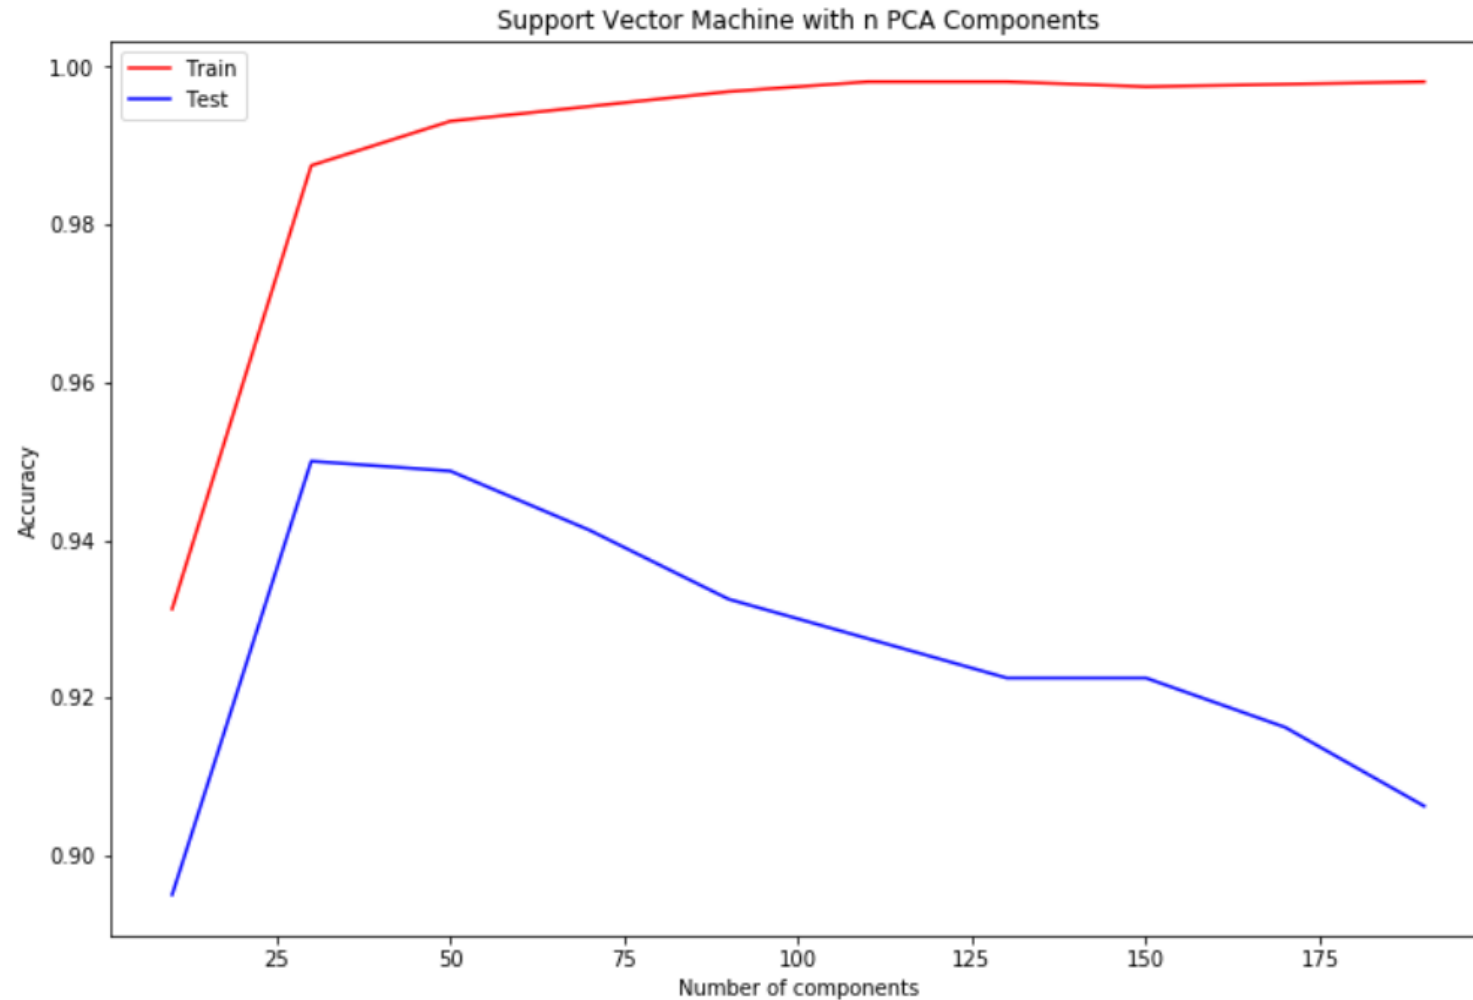


Figure 2-1. Trade-off of model complexity against training and test accuracy

# Complexity vs. Accuracy



# Supervised Learning Algorithms Overview



- **For each supervised learning method we'll explore:**
  - *How the method works conceptually at a high level.*
  - *What kind of feature preprocessing is typically needed.*
  - *Key parameters that control model complexity, to avoid under-and over-fitting.*
  - *Positives and negatives of the learning method.*
- **This week:**
  - *K-nearest neighbors*
  - *Linear model fit using least-squares*

# *K-nearest Neighbors*






Figure 2-4. Predictions made by the one-nearest-neighbor model on the forge dataset

- To predict, find the closest data points in the training dataset—its “nearest neighbors.”
- Import parameters:
  - Model complexity
    - `n_neighbors`
  - Model fitting
    - Metric: distance between data points
      - Default: Euclidean



# Less Neighbors, More Complex



# Pros and Cons



## Pros

- Easy to understand
- Easy to interpret
- Fast training

## Cons

- With large training dataset, prediction can be very slow
- Does not perform well on many features
- Does not perform well on sparse data (data with majority 0s)

# Applications of KNN in Real World



- Useful in comparing similarities.
  - Concept search
  - Recommender systems
- Use as a benchmark for other more advanced algorithms
- Find out outliers

# Linear Models



- $y = b + wX$
- $y$  is the label, or target to be predicted, or dependent variable
- $b$  is the intercept
- $w$  is a weight vector, or coefficient vector
- $X$  is the feature matrix, or predictors matrix, or independent
- $X = (x_0, x_1, x_2, x_3, \dots x_{n-1}, x_n)$ ,  $w = (w_0, w_1, w_2, w_3, \dots w_{n-1}, w_n)$
- $\hat{y} = \hat{b} + \hat{w}_0x_0 + \hat{w}_1x_1 + \dots + \hat{w}_{n-1}x_{n-1} + \hat{w}_nx_n$
- $\hat{b}, \hat{w}$  are the parameters to be estimated

# Linear Models




- Regressions:
  - Linear regression (ordinary least squares)
  - Ridge regression
  - Lasso regression
  - ElasticNet regression
  - SGD regression (large numbers of samples and features)
- Classifications:
  - LinearSVC
  - Logistic regression
  - ElasticNet classifier
  - SGD classifier



# Linear Regression (Ordinary Least Squares)



- Finds  $w$  and  $b$  that minimizes the sum of squared differences(RSS) over the training data between predicted target and actual target values.
- a.k.a. mean squared error of the linear model
- No parameters to control model complexity



$$RSS(w, b) = \sum_{i=1}^N (y_i - (w \cdot x_i + b))^2$$

**Loss function**

# Linear Regression (Ordinary Least Squares)



- **Parameters are estimated from training data.**
- **There are many different ways to estimate  $w$  and  $b$ :**
  - *Different methods correspond to different "fit" criteria and goals and ways to control model complexity.*
- **The learning algorithm finds the parameters that optimize an objective function, typically to minimize some kind of loss function of the predicted target values vs. actual target values.**

## Fit linear regression model

```
clf=LinearRegression()  
clf.fit(X_train,y_train)
```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

## Check the R square of the regression model on training data

```
clf.score(X_train,y_train)
```

```
0.77297187265711575
```

## Check the R square of the regression model on test data

```
clf.score(X_test,y_test)
```

```
0.58920115191864408
```

$$y = b + wX$$
$$y = 38.14 - 0.1184X_0 + 0.0448X_1 + 0.0059X_2 + \dots - 0.4936X_{12}$$

## Display the intercept and coefficients

```
print('The intercept is {0} \nand coefficient vector is \n{1}'.format(clf.intercept_,clf.coef_))
```

```
The intercept is [ 38.13869271]
```

```
and coefficient vector is
```

```
[[ -1.18410318e-01  4.47550643e-02  5.85674689e-03  2.34230117e+00  
 -1.61634024e+01  3.70135143e+00 -3.04553661e-03 -1.38664542e+00  
  2.43784171e-01 -1.09856157e-02 -1.04699133e+00  8.22014729e-03  
 -4.93642452e-01]]
```



# Ridge Regression



Ridge regression learns  $w$ ,  $b$  using the same least-squares criterion but adds a penalty for large variations in  $w$  parameters

$$RSS_{RIDGE}(w, b) = \sum_{i=1}^N (y_i - (w \cdot x_i + b))^2 + \alpha \sum_{j=1}^p w_j^2$$

Once the parameters are learned, the ridge regression prediction formula is the same as ordinary least-squares.

The addition of a parameter penalty is called regularization. Regularization prevents overfitting by restricting the model, typically to reduce its complexity.

Ridge regression uses L2 regularization: minimize sum of squares of  $w$  entries

The influence of the regularization term is controlled by the  $\alpha$  parameter.

Higher alpha means more regularization and simpler models.

# Lasso Regression



Lasso regression is another regularized linear regression that uses an **L1** regularization penalty for training (instead of ridge's **L2** penalty)

**L1 penalty:** Minimize the sum of the absolute values of the coefficients

$$RSS_{LASSO}(\mathbf{w}, b) = \sum_{\{i=1\}}^N (y_i - (\mathbf{w} \cdot \mathbf{x}_i + b))^2 + \alpha \sum_{\{j=1\}}^p |w_j|$$

This has the effect of setting parameter weights in  $\mathbf{w}$  to zero for the least influential variables. This is called a sparse solution: a kind of feature selection

The parameter  $\alpha$  controls amount of L1 regularization (default = 1.0).

The prediction formula is the same as ordinary least-squares.

**When to use ridge vs lasso regression:**

- *Many small/medium sized effects: use ridge.*
- *Only a few variables with medium/large effect: use lasso.*



# ElasticNet Regression



- Linear regression with combined L1 and L2 priors as regularizer.

$$\min_w \frac{1}{2n_{\text{samples}}} \|Xw - y\|_2^2 + \alpha \rho \|w\|_1 + \frac{\alpha(1 - \rho)}{2} \|w\|_2^2$$

- This combination allows for learning a sparse model where few of the weights are non-zero like Lasso, while still maintaining the regularization properties of Ridge.
- Parameters:
  - Alpha: High  $\alpha$  means more regularization.
  - l1\_ratio: with  $0 \leq \text{l1\_ratio} \leq 1$ . For  $\text{l1\_ratio} = 0$  the penalty is an L2 penalty. For  $\text{l1\_ratio} = 1$  it is an L1 penalty. For  $0 < \text{l1\_ratio} < 1$ , the penalty is a combination of L1 and L2.

# Polynomial Features with Linear Regression



$$\mathbf{x} = (x_0, x_1) \longrightarrow \mathbf{x}' = (x_0, x_1, x_0^2, x_0x_1, x_1^2)$$

$$\hat{y} = \hat{w}_0x_0 + \hat{w}_1x_1 + \hat{w}_{00}x_0^2 + \hat{w}_{01}x_0x_1 + \hat{w}_{11}x_1^2 + b$$

**Generate new features consisting of all polynomial combinations of the original two features  $(x_0, x_1)$ .**


**The *degree* of the polynomial specifies how many variables participate at a time in each new feature (above example: degree 2)**

**This is still a weighted linear combination of features, so it's still a linear model, and can use same least-squares estimation method for  $w$  and  $b$ .**


# Polynomial Features with Linear Regression



Complex regression problem with one input variable



Complex regression problem with one input variable



# Polynomial Features with Linear Regression



## **Why would we want to transform our data this way?**

- *To capture interactions between the original features by adding them as features to the linear model.*
- *To make a classification problem easier (we'll see this later).*

## **More generally, we can apply other non-linear transformations to create new features**

- *(Technically, these are called non-linear basis functions)*


## **Beware of polynomial feature expansion with high as this can lead to complex models that overfit**

- *Thus, polynomial feature expansion is often combined with a regularized learning method like ridge regression.*

# Linear Regression





Input features



$$\hat{y} = \hat{b} + \hat{w}_1 \cdot x_1 + \cdots \hat{w}_n \cdot x_n$$



# Logistic Regression



The logistic function transforms real-valued input to an output number  $y$  between 0 and 1, interpreted as the probability the input object belongs to the positive class, given its input features  $(x_0, x_1, \dots, x_n)$

$$\begin{aligned}\hat{y} &= \text{logistic}(\hat{b} + \hat{w}_1 \cdot x_1 + \dots \hat{w}_n \cdot x_n) \\ &= \frac{1}{1 + \exp[-(\hat{b} + \hat{w}_1 \cdot x_1 + \dots \hat{w}_n \cdot x_n)]}\end{aligned}$$


# Regularization of Logistic Regression



**L2 regularization is 'on' by default (like ridge regression)**

**Parameter C controls amount of regularization (default 1.0)**

**As with regularized linear regression, it can be important to normalize all features so that they are on the same scale.**



# Linear Support Vector Machines Classifier



**The strength of regularization is determined by  $C$**

**Larger values of  $C$ : less regularization**

- *Fit the training data as well as possible*
- *Each individual data point is important to classify correctly*

**Smaller values of  $C$ : more regularization**

- *More tolerant of errors on individual data points*

# Important Parameters



## **Model complexity**

- **alpha:** weight given to the L1 or L2 regularization term in regression models
  - *default = 1.0*
- **C:** regularization weight for LinearSVC and LogisticRegression classification models
  - *default = 1.0*

# The Need for Feature Normalization



- **Important for some machine learning methods that all features are on the same scale (e.g. faster convergence in learning, more uniform or 'fair' influence for all weights)**
  - *e.g. regularized regression, k-NN, support vector machines, neural networks*

# Linear Models



## **Pros:**

- Simple and easy to train.
- Fast prediction.
- Scales well to very large datasets.
- Works well with sparse data.
- Reasons for prediction are relatively easy to interpret.

## **Cons:**

- For lower-dimensional data, other models may have superior generalization performance.
- For classification, data may not be linearly separable (more on this in SVMs with non-linear kernels)

# Questions?

